




Boosting

Adaboost, Gradient Boost, XGBoost

Seunghan Lee

20.03.18 (Wed)



Contents

1

What is Boosting?

1) Ensemble method / 2) Bagging vs Boosting

2

1. Adaboost

2. Gradient Boost

3. XGBoost

3

Boosting with Python



1. What is Boosting?

- 1) Ensemble method
- 2) Bagging vs Boosting

1. What is Boosting?

(1) Ensemble

*In statistics and machine learning, **ensemble methods** use **multiple learning algorithms** to obtain better predictive performance than could be obtained from any of the constituent learning algorithms alone. (Wikipedia)*

- Use “**multiple models**” to predict the output
(both regression & classification)
- Several Weak learner > One Strong learner
- ex) Bagging & Boosting

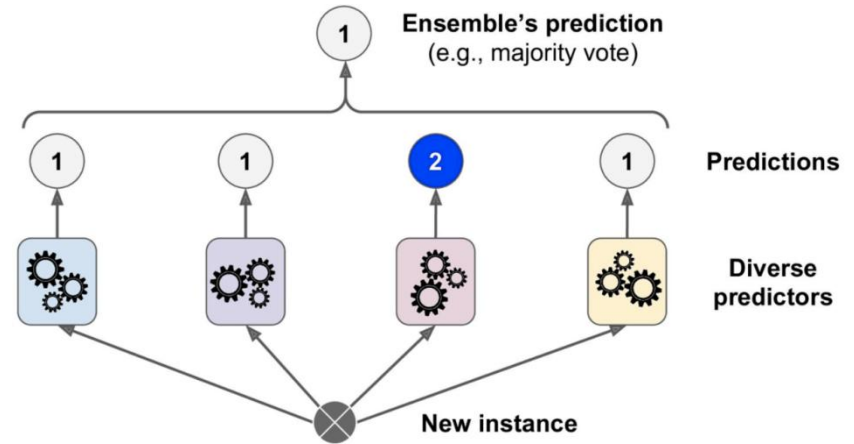


Figure 7-2. Hard voting classifier predictions

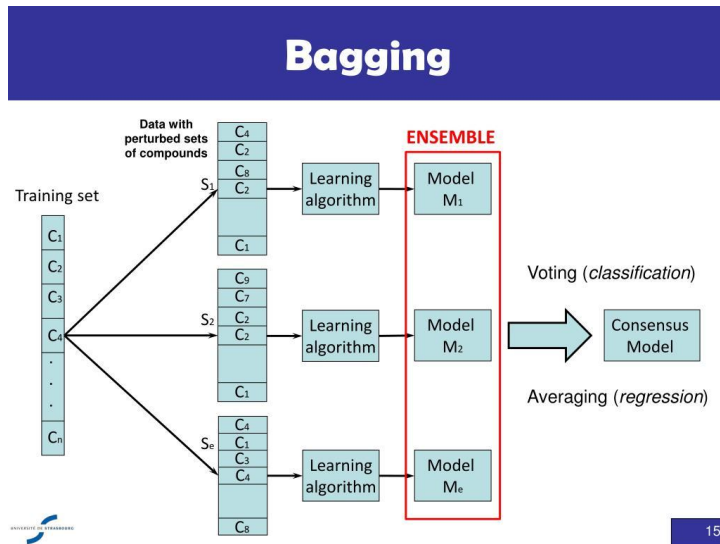
1. What is Boosting?

(2) Bagging vs Boosting

Steps of Bagging

- 1) Random sampling with replacement (= Bootstrapping)
- 2) Select subset of features randomly
- 3) Grow the tree to the largest
- 4) Repeat 1)~3) N times (total of N models)
-> make prediction based on these N models

<https://www.slideserve.com/tarala/short-overview-of-weka>

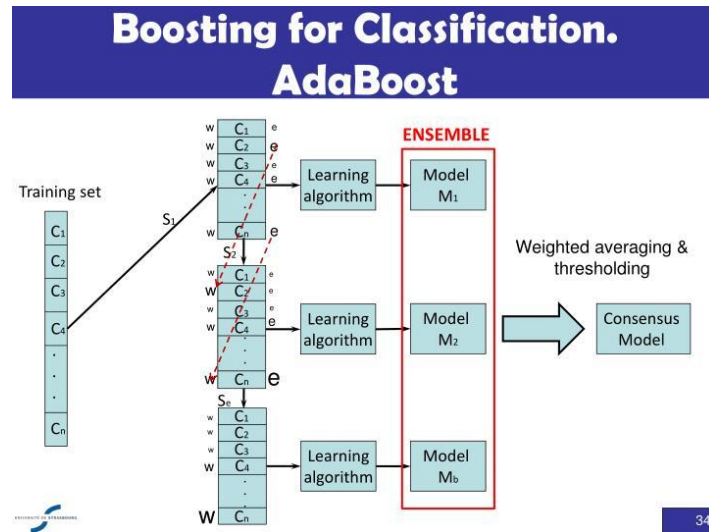


1. What is Boosting?

(2) Bagging vs Boosting

Steps of Boosting

- 1) First random sample (d_1) without replacement
-> train a weak learner C_1
- 2) Second random sample (d_2) without replacement, considering the misclassified sample in learner C_1
-> train a weak learner C_2
- 3) Third random sample (d_3) without replacement, considering the misclassified sample in learner C_1, C_2
-> train a weak learner C_3
- 4) Combine all the weak learners ($C_1 \sim C_n$) via majority voting.



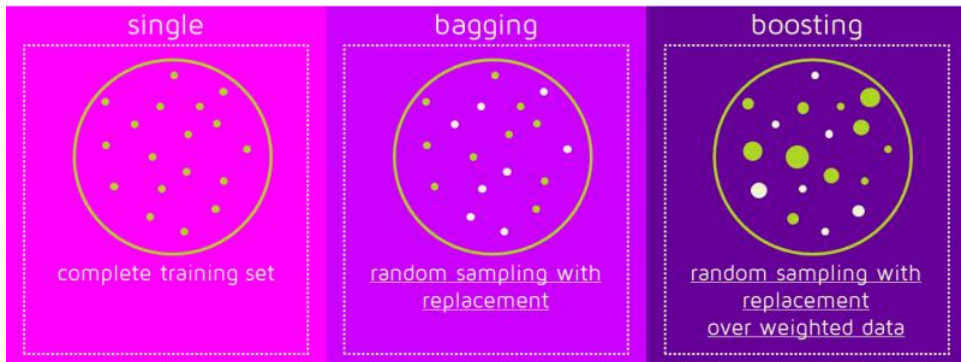
1. What is Boosting?

(2) Bagging vs Boosting

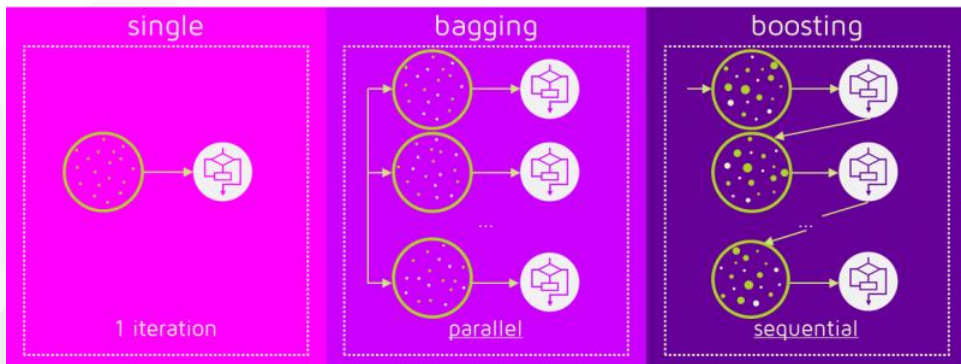
	Bagging	Boosting
	Parallel Ensemble (Each model is independent)	Sequential Ensemble (Next Model is dependent to the previous model)
Sampling	Random sample (Bootstrapping)	Higher vote to misclassified sample
Goal	Reduce Variance	Reduce Bias
Base Learner	Unstable Learner (ex. Decision Tree)	Stable Learner (ex. Stump)
Example)	Random Forest	XGBoost, LightGBM

1. What is Boosting?

(2) Bagging vs Boosting



way of sampling



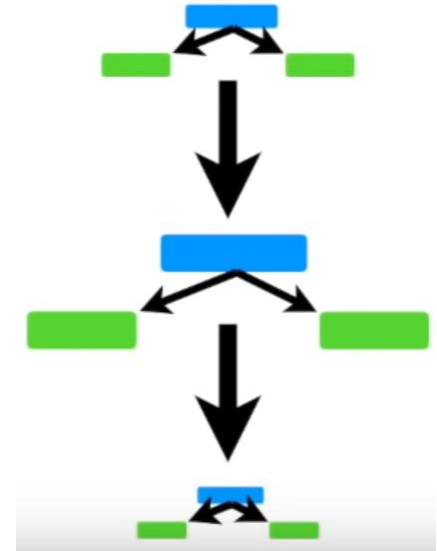
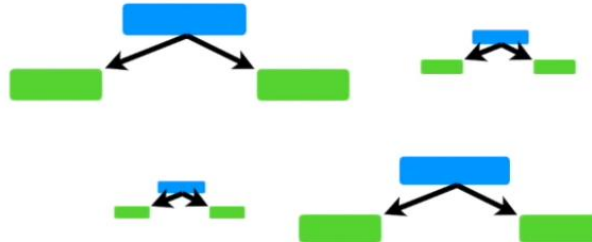
way of making models



2-a. Adaboost

What is Adaboost?

- Adaboost = **Adaptive Boosting**
(Adaptive? “more likely to choose the misclassified samples” from the previous learners)
- base learner : **weak learners** (not very accurate)
ex) **Stump** (= a tree with 1 node & 2 leaves)
- **weighted sum** of the output of weak learners
(= unlike RF, learners get different vote to the final output!)
- weak learners are **dependent** to each other! (made **sequentially**)



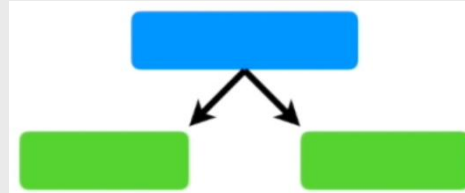
2-a. Adaboost

How does Adaboost work?

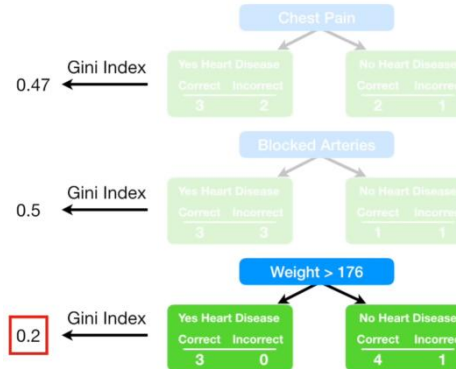
Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8
No	Yes	156	No	1/8
No	Yes	125	No	1/8
Yes	No	168	No	1/8
Yes	Yes	172	No	1/8

(How each data is important to be correctly classified) -> will be updated!

1) Make a Stump



with the variable which makes the best classification



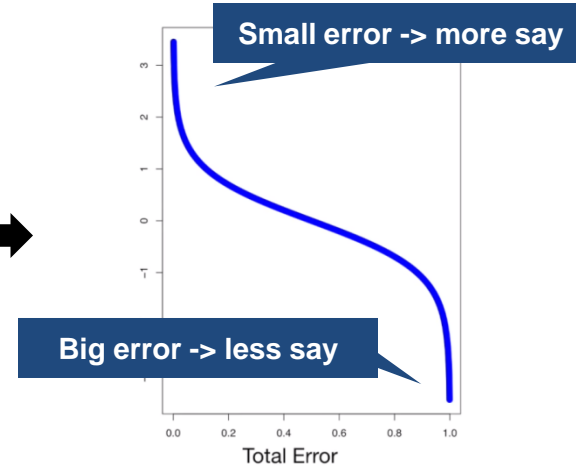
In this example,
“weight > 176” will be chosen

2-a. Adaboost

How does Adaboost work?

“How much say” does each stump has?

$$\text{Amount of Say} = \frac{1}{2} \log\left(\frac{1 - \text{Total Error}}{\text{Total Error}}\right)$$



Example)



Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8




$$\text{Amount of Say} = \frac{1}{2} \log(7) = 0.97$$

2-a. Adaboost

How does Adaboost work?

“How much say” does each stump has?


Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8



Need to update!
(to a SMALLER value)



Need to update!
(to a LARGER value)



Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8



$$\text{Amount of Say} = \frac{1}{2} \log(7) = 0.97$$

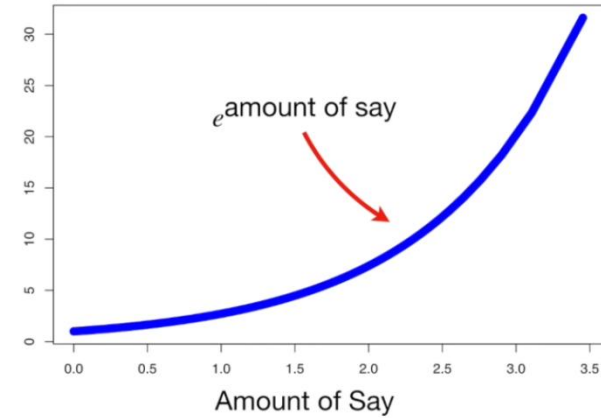
2-a. Adaboost

How does Adaboost work?

< for the “Misclassified” sample >

$$\text{New Sample Weight} = \text{sample weight} \times e^{\text{amount of say}}$$

(increase the weight for the “incorrectly classified” sample!)



Example)

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8

$$\text{Amount of Say} = \frac{1}{2} \log(7) = 0.97$$

$$\frac{1}{8} e^{0.97} = \frac{1}{8} \times 2.64 = 0.33$$

Before) 0.125

After) **0.33**


2-a. Adaboost

How does Adaboost work?

< for the “Correctly classified” sample >

$$\text{New Sample Weight} = \text{sample weight} \times e^{-\text{amount of say}}$$

(decrease the weight for the “correctly classified” sample!)



Example)

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	1/8
No	Yes	180	Yes	1/8
Yes	No	210	Yes	1/8
Yes	Yes	167	Yes	1/8

$$\text{Amount of Say} = \frac{1}{2} \log(7) = 0.97$$

$$\frac{1}{8} e^{-0.97} = \frac{1}{8} \times 0.38 = 0.05$$

Before) 0.125

After) **0.05**

2-a. Adaboost

How does Adaboost work?

After updating the weight
(+ normalization)

New Weight	Norm. Weight
0.05	0.07
0.05	0.07
0.05	0.07
0.33	0.49
0.05	0.07
0.05	0.07
0.05	0.07
0.05	0.07



This becomes the new sample weight!

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	0.07
No	Yes	180	Yes	0.07
Yes	No	210	Yes	0.07
Yes	Yes	167	Yes	0.49
No	Yes	156	No	0.07
No	Yes	125	No	0.07
Yes	No	168	No	0.07
Yes	Yes	172	No	0.07

2-a. Adaboost

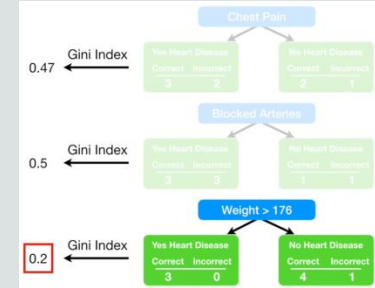
How does Adaboost work?

Make the next stump,
with the “new sample weight”

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	0.07
No	Yes	180	Yes	0.07
Yes	No	210	Yes	0.07
Yes	Yes	167	Yes	0.49
No	Yes	156	No	0.07
No	Yes	125	No	0.07
Yes	No	168	No	0.07
Yes	Yes	172	No	0.07

Option 1)

Use “weighted Gini index” to select
the feature to make new stump



Option 2)

Duplicate copies of the samples
(bigger ‘sample weight’ ->
more likely to be selected)

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease
No	Yes	156	No
Yes	Yes	167	Yes
No	Yes	125	No
Yes	Yes	167	Yes
Yes	Yes	167	Yes
Yes	Yes	172	No
Yes	Yes	205	Yes
Yes	Yes	167	Yes

2-a. Adaboost

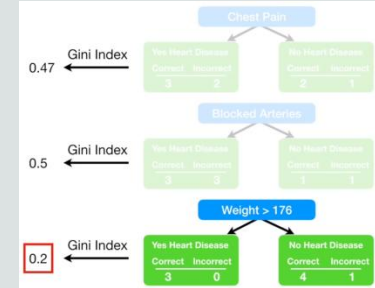
How does Adaboost work?

Make the next stump,
with the “new sample weight”

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
Yes	Yes	205	Yes	0.07
No	Yes	180	Yes	0.07
Yes	No	210	Yes	0.07
Yes	Yes	167	Yes	0.49
No	Yes	156	No	0.07
No	Yes	125	No	0.07
Yes	No	168	No	0.07
Yes	Yes	172	No	0.07

Option 1)

Use “weighted Gini index” to select
the feature to make new stump



Option 2)

Duplicate copies of the samples
(bigger ‘sample weight’ ->
more likely to be selected)

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease
No	Yes	156	No
Yes	Yes	167	Yes
No	Yes	125	No
Yes	Yes	167	Yes
Yes	Yes	167	Yes
Yes	Yes	172	No
Yes	Yes	205	Yes
Yes	Yes	167	Yes

2-a. Adaboost

How does Adaboost work?

Chest Pain	Blocked Arteries	Patient Weight	Heart Disease	Sample Weight
No	Yes	156	No	1/8
Yes	Yes	167	Yes	1/8
No	Yes	125	No	1/8
Yes	Yes	167	Yes	1/8
Yes	Yes	167	Yes	1/8
Yes	Yes	172	No	1/8
Yes	Yes	205	Yes	1/8
Yes	Yes	167	Yes	1/8

Give equal weight to all the samples!

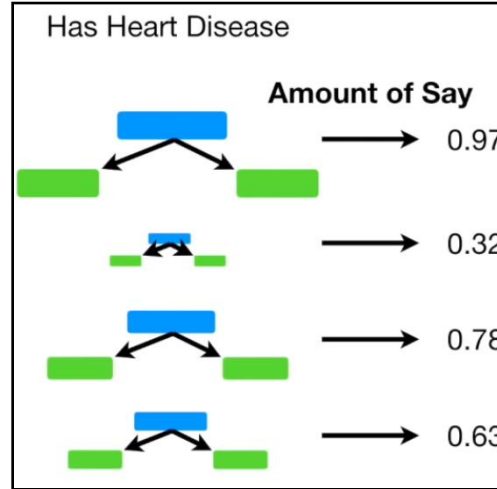
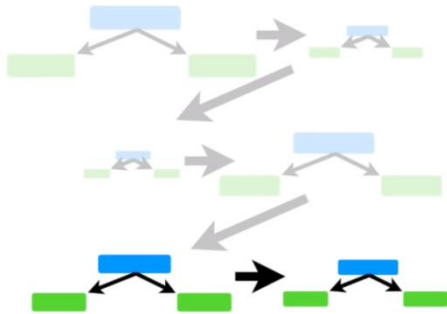
Same sample weights doesn't mean same importance to every sample!

(important samples are duplicated!)

2-a. Adaboost

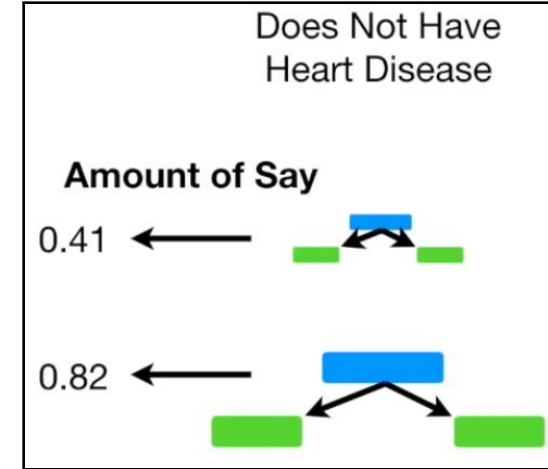
How does Adaboost work?

Make models(stumps) sequentially



Total Sum : 2.7

>



Total Sum : 1.23

How to predict
the final output?

Result : "Has Heart Disease"

How does Adaboost work?

Binary AdaBoost

Input:

- Data: $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^N$, where $x^{(i)} \in \mathcal{X}$ and $y^{(i)} \in \mathcal{Y} = \{-1, 1\}$ for $i = 1, \dots, N$
- Total number of rounds: T

Initialize: Define the initial probability distribution $D_1(i) = \frac{1}{N}$ for $i = 1, \dots, N$.

Do for $t = 1, \dots, T$:

- Train using the probability distribution D_t on $\{1, \dots, N\}$.
- Get a hypothesis (classifier) $h_t : \mathcal{X} \rightarrow \mathcal{Y}$ that has a training error rate ε_t (with respect to D_t) given by

$$\varepsilon_t = \sum_{i: h_t(x^{(i)}) \neq y^{(i)}} D_t(i).$$

- If $\varepsilon_t > \frac{1}{2}$, then set $T = t - 1$ and abort the loop.
- Set $\alpha_t = \frac{1}{2} \log \left(\frac{1 - \varepsilon_t}{\varepsilon_t} \right)$.
- Define the new probability distribution D_{t+1} by setting

$$\begin{aligned} D_{t+1}(i) &= \frac{D_t(i)}{Z_t} \cdot \begin{cases} e^{-\alpha_t} & \text{if } y^{(i)} = h_t(x^{(i)}) \\ e^{\alpha_t} & \text{if } y^{(i)} \neq h_t(x^{(i)}) \end{cases} \\ &= \frac{D_t(i)}{Z_t} \exp(-\alpha_t y^{(i)} h_t(x^{(i)})), \end{aligned}$$

where

$$Z_t = \sum_{i=1}^N D_t(i) \exp(-\alpha_t y^{(i)} h_t(x^{(i)})).$$

Output: Define $g(x) = \sum_{t=1}^T \alpha_t h_t(x)$. The final hypothesis (classifier) $H(x)$ is

$$H(x) = \text{sgn}(g(x)).$$

binary



Extend to multi-class

AdaBoost.M1

Input:

- Data: $\mathcal{D} = \{(x^{(i)}, y^{(i)})\}_{i=1}^N$, where $x^{(i)} \in \mathcal{X}$ and $y^{(i)} \in \mathcal{Y}$ for $i = 1, \dots, N$. Here \mathcal{Y} is a set of K elements, which, as usual, is identified with $\{1, \dots, K\}$
- Total number of rounds: T

Initialize: Define the initial probability distribution $D_1(i) = \frac{1}{N}$ for $i = 1, \dots, N$.

Do for $t = 1, 2, \dots, T$:

- Train using the probability distribution D_t on $\{1, \dots, N\}$.
- Get a hypothesis (classifier) $h_t : \mathcal{X} \rightarrow \mathcal{Y}$ that has a training error rate ε_t (with respect to D_t) given by

$$\varepsilon_t = \sum_{i: h_t(x^{(i)}) \neq y^{(i)}} D_t(i).$$

- If $\varepsilon_t > \frac{1}{2}$, then set $T = t - 1$ and abort the loop.
- Set $\alpha_t = \frac{1}{2} \log \left(\frac{1 - \varepsilon_t}{\varepsilon_t} \right)$.
- Define the new probability distribution D_{t+1} by setting

$$\begin{aligned} D_{t+1}(i) &= \frac{D_t(i)}{Z_t} \cdot \begin{cases} e^{-\alpha_t} & \text{if } y^{(i)} = h_t(x^{(i)}) \\ e^{\alpha_t} & \text{if } y^{(i)} \neq h_t(x^{(i)}) \end{cases} \\ &= \frac{D_t(i)}{Z_t} \exp(-\alpha_t S(h_t(x^{(i)}), y^{(i)})), \end{aligned}$$

where $Z_t = \sum_{i=1}^N D_t(i) \exp(-\alpha_t S(h_t(x^{(i)}), y^{(i)}))$, and $S(a, b) = \mathbb{I}(a = b) - \mathbb{I}(a \neq b)$ so that $S(a, b) = 1$ if $a = b$ and $S(a, b) = -1$ if $a \neq b$.

Output: the final hypothesis (classifier)

$$H(x) = \argmax_{y \in \mathcal{Y}} \sum_{t: h_t(x) = y} \alpha_t.$$

Multi-class



Adaboost Summary

- 1) Combine **weak learners** to make classifications
- 2) Some learners **have more say** than over learners
- 3) Each learner is made by considering the result of the **previous learner's mistake**



2-b. Gradient Boost

2-b. Gradient Boost

What is Gradient Boost?

- not so much different with Adaboost

- difference?

- 1) Combine **weak learners** to make classifications
- 2) Some learners **have more say** than over learners
- 3) Each learner is made by considering the result of the **previous learner's mistake**

Can be expressed in more general terms! (by using “**gradient**”)

+ use trees larger than stumps as base learners

Minimizing the cost(loss) function

1. Gradient Boosting Regression

2. Gradient Boosting Classification

$$\begin{aligned} y &= f(x) + \underline{e1} \\ e1 &= g(x) + \underline{e2} \\ e2 &= h(x) + e3 \end{aligned}$$

$$Y = f(x) + g(x) + h(x) + \dots$$

1. Gradient Boost Regression

Height (m)	Favorite Color	Gender	Weight (kg)
1.6	Blue	Male	88
1.6	Green	Female	76
1.5	Blue	Female	56
1.8	Red	Male	73
1.5	Green	Male	77
1.4	Blue	Female	57

1

First, make a “leaf (= guess for the y value)”

(initial leaf : average of all the y values)

Average Weight

71.2

(= first predicted y)

1. Gradient Boost Regression

Height (m)	Favorite Color	Gender	Weight (kg)
1.6	Blue	Male	88
1.6	Green	Female	76
1.5	Blue	Female	56
1.8	Red	Male	73
1.5	Green	Male	77
1.4	Blue	Female	57

1

First, make a “**leaf** (= guess for the y value)”

(initial leaf : average of all the y values)

Average Weight
71.2

(= first predicted y)

Height (m)	Favorite Color	Gender	Weight (kg)	Residual
1.6	Blue	Male	88	16.8
1.6	Green	Female	76	4.8
1.5	Blue	Female	56	-15.2
1.8	Red	Male	73	1.8
1.5	Green	Male	77	5.8
1.4	Blue	Female	57	-14.2

2

Then, predict the “**residual**”

(build the next model (tree) predicting the residuals)

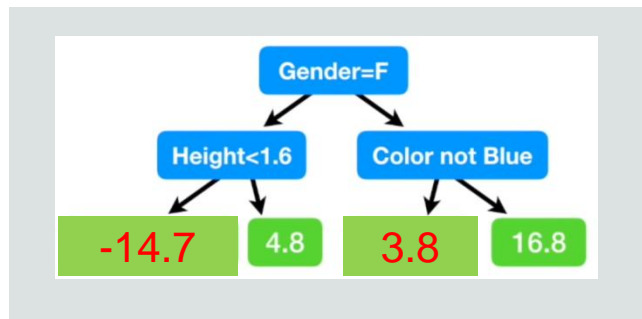
1. Gradient Boost Regression

Height (m)	Favorite Color	Gender	Weight (kg)	Residual
1.6	Blue	Male	88	16.8
1.6	Green	Female	76	4.8
1.5	Blue	Female	56	-15.2
1.8	Red	Male	73	1.8
1.5	Green	Male	77	5.8
1.4	Blue	Female	57	-14.2

(usually 8~32 leaves)



Replace the residuals into their average




2-b. Gradient Boost

1. Gradient Boost Regression

Average Weight

71.2

+



Height (m)	Favorite Color	Gender	Weight (kg)
1.6	Blue	Male	88

100% correct (no error)

Overfitting!


2-b. Gradient Boost

1. Gradient Boost Regression

Average Weight

71.2

+



Height (m)	Favorite Color	Gender	Weight (kg)
1.6	Blue	Male	88

100% correct (no error)


Overfitting!

Average Weight

71.2

+

$lr \times$



Height (m)	Favorite Color	Gender	Weight (kg)
1.6	Blue	Male	72.9

(If learning rate = 0.1)

$$71.2 + (0.1 \times 16.8) = 72.9$$



1. Gradient Boost Regression


new residuals
(becomes smaller)

Height (m)	Favorite Color	Gender	Weight (kg)	Residual
1.6	$88 - (71.2 + 0.1 \times 16.8) =$			15.1
1.6	Green	Female	76	4.3
1.5	Blue	Female	56	-13.7
1.8	Red	Male	73	1.4
1.5	Green	Male	77	5.4
1.4	Blue	Female	57	-12.7

1. Gradient Boost Regression

new residuals
(becomes smaller)

Height (m)	Favorite Color	Gender	Weight (kg)	Residual
1.6				$88 - (71.2 + 0.1 \times 16.8) = 15.1$
1.6	Green	Female	76	4.3
1.5	Blue	Female	56	-13.7
1.8	Red	Male	73	1.4
1.5	Green	Male	77	5.4
1.4	Blue	Female	57	-12.7



3

Then, predict the “**residual’s residual**”
(build the next model (tree) predicting the residuals)

1. Gradient Boost Regression

Height (m)	Favorite Color	Gender	Weight (kg)
1.6	Blue	Male	74.4

$$71.2 + (0.1 \times 16.8) + (0.1 \times 15.1) = 74.4$$

$$y = f(x) + e1$$

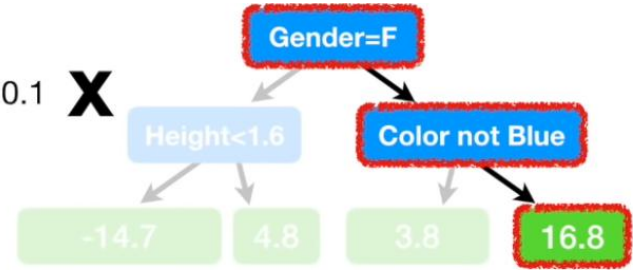
$$e1 = g(x) + e2$$

$$e2 = h(x) + e3$$

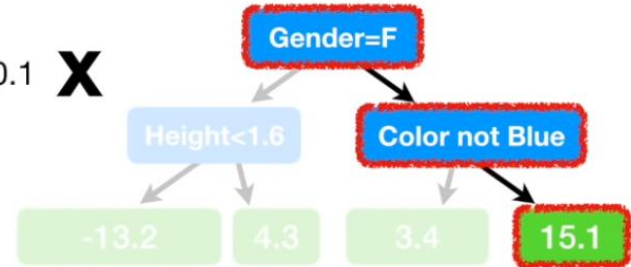
Average Weight

71.2


+ 0.1 X



+ 0.1 X



1. Gradient Boost Regression



Keep making the **sequential** trees!

(until specified level, or until the residuals
doesn't significantly drops)

1. Gradient Boost Regression

Mathematical Expression

$$\text{Loss Function} = \frac{1}{2} (\text{Observed} - \text{Predicted})^2$$



To minimize this loss function

$$\frac{d}{d \text{ Predicted}} \frac{1}{2} (\text{Observed} - \text{Predicted})^2 = -(\text{Observed} - \text{Predicted}) = 0$$

Height (m)	Favorite Color	Gender	Weight (kg)
1.6	Blue	Male	88
1.6	Green	Female	76
1.5	Blue	Female	56
1.8	Red	Male	73
1.5	Green	Male	77
1.4	Blue	Female	57

$$\begin{aligned} & \frac{1}{2} (88 - \text{Predicted})^2 + \rightarrow -(88 - \text{Predicted}) + \\ & \frac{1}{2} (76 - \text{Predicted})^2 + \rightarrow -(76 - \text{Predicted}) + \\ & \frac{1}{2} (56 - \text{Predicted})^2 \rightarrow -(56 - \text{Predicted}) \\ & \frac{d}{d \text{ Predicted}} \end{aligned} = 0$$

1. Gradient Boost Regression

Mathematical Expression

$$\text{Loss Function} = \frac{1}{2} (\text{Observed} - \text{Predicted})^2$$



To minimize this loss function

$$\frac{d}{d \text{ Predicted}} \frac{1}{2} (\text{Observed} - \text{Predicted})^2 = -(\text{Observed} - \text{Predicted}) = 0$$

Height (m)	Favorite Color	Gender	Weight (kg)
1.6	Blue	Male	88
1.6	Green	Female	76
1.5	Blue	Female	56
1.8	Red	Male	75
1.5	Green	Male	77
1.4	Blue	Female	57

That is why we kept making
the predicted value as the “**AVERAGE**” value!

1. Gradient Boost Regression

<https://ascelibrary.org/cms/asset/>

Gradient Boosted Regression Tree

1. **Input:** $D = \{(x_1, y_1), \dots, (x_N, y_N)\}, \theta, \gamma$
 2. **Output:** $F(x) = \sum_{i=0}^M F_i(x)$
 3. Initialize $F_0(x) = \arg \min_{\beta} \sum_{i=0}^N L(y_i, \beta)$
 4. **While** ($m < M$)
 5. $d_i = -[\partial L(y_i, F(x_i)) / \partial F(x_i)]_{F(x_i)=F_{m-1}(x_i)}$
 6. $\vartheta = \{(x_i, d_i)\}, i = 1, N$
 7. $g(x) = \text{FITREGRETREE}(\vartheta, \theta)$
 8. $\rho_m = \arg \min_{\rho} \sum_{i=1}^N L(y_i, F_{m-1}(x) + \rho g(x))$
 9. $F_m(x) = F_{m-1}(x) + \gamma \rho_m g(x)$
 13. **End while**
-

Done!

2-b. Gradient Boost

1. Gradient Boost Regression

Gradient Boosted Regression Tree

1. **Input:** $D = \{(x_1, y_1), \dots, (x_N, y_N)\}, \theta, \gamma$
2. **Output:** $F(x) = \sum_{i=0}^M F_i(x)$
3. Initialize $F_0(x) = \arg \min_{\beta} \sum_{i=0}^N L(y_i, \beta)$
4. **While** ($m < M$)
 - Loss Function : MSE
 - $d_i = -[\partial L(y_i, F(x_i)) / \partial F(x_i)]_{F(x_i)=F_{m-1}(x_i)}$
 - $\vartheta = \{(x_i, d_i)\}, i = 1, N$
 - $g(x) = \text{FITREGRETREE}(\vartheta, \theta)$
 - $\rho_m = \arg \min_{\rho} \sum_{i=1}^N L(y_i, F_{m-1}(x) + \rho g(x))$
 - $F_m(x) = F_{m-1}(x) + \gamma \rho_m g(x)$
13. **End while**

$$y = f(x) + e_1$$
$$e_1 = g(x) + e_2$$
$$e_2 = h(x) + e_3$$

$$Y = f(x) + g(x) + h(x) + \dots$$

First prediction = “average” of all y values

Height (m)	Color	Gender	Weight (kg)	Residual
1.6	Blue	Male	80	16.8
1.6	Green	Female	76	4.8
1.5	Blue	Female	66	-15.2
1.8	Red	Male	73	1.8
1.5	Green	Male	77	5.8
1.4	Blue	Female	57	-14.2



$$Y = f(x) + g(x) + e_2$$



2. Gradient Boost Classification

Likes Popcorn	Age	Favorite Color	Loves Troll 2
Yes	12	Blue	Yes
Yes	87	Green	Yes
No	44	Blue	No
Yes	19	Red	No
No	32	Green	Yes
No	14	Blue	Yes



4 Yes, 2 No

Odds ratio (of yes): $4/2 = 2$

Odds ratio (of no) : $2/4 = 0.5$

$\text{Log}(2) = 0.7$

$\text{Log}(0.5) = -0.7$

Log Odds Ratio?

1

First, make a “leaf (= guess for the y value)”

(initial leaf : log odds ratio of the highest value)

$\text{log}(4/2) = 0.7$

(= first predicted y)



2. Gradient Boost Classification

1) Log Odds Ratio (= Logit)

$$\text{logit}(p) = \log\left(\frac{p}{1-p}\right)$$

If p is a probability, then $p/(1 - p)$ is the corresponding odds

2. Gradient Boost Classification

1) Log Odds Ratio (= Logit)

$$\text{logit}(p) = \log\left(\frac{p}{1-p}\right)$$

If p is a probability, then $p/(1-p)$ is the corresponding odds

Example)

Chance of winning?

DSL	YBIGTA
5	2



$$\frac{5}{2}$$



$$(5/7) / (2/7) = 5/2$$

Odds are 2.5
that DSL will win!
(probability (p) = $5/7$)



$$\log(5/2) = 0.916$$

Why log?





2. Gradient Boost Classification

2) Logistic Function (= Sigmoid)

$$f(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1}$$

As a method to turn 'log odds ratio' back to 'probability'



2. Gradient Boost Classification

2) Logistic Function (= Sigmoid)

$$f(x) = \frac{1}{1 + e^{-x}} = \frac{e^x}{e^x + 1}$$

As a method to turn 'log odds ratio' back to 'probability'

Example)

(The probability was 5/7)

$$\log(5/2) = 0.916 \quad \longrightarrow \quad f(\log(5/2)) = 5/7$$

```
>>> 5/7
0.7142857142857143
>>>
>>> 1/(1+np.exp(-np.log(5/2)))
0.7142857142857143
```



2. Gradient Boost Classification

Likes Popcorn	Age	Favorite Color	Loves Troll 2
Yes	12	Blue	Yes
Yes	87	Green	Yes
No	44	Blue	No
Yes	19	Red	No
No	32	Green	Yes
No	14	Blue	Yes



4 Yes, 2 No

Odds ratio (of yes): $4/2 = 2$

Odds ratio (of no) : $2/4 = 0.5$

$\text{Log}(2) = 0.7$

$\text{Log}(0.5) = -0.7$

1

First, make a “leaf (= guess for the y value)”

(initial leaf : log odds ratio of the highest value)

$\text{log}(4/2) = 0.7$

(= first predicted y)

2. Gradient Boost Classification

Likes Popcorn	Age	Favorite Color	Loves Troll 2
Yes	12	Blue	Yes
Yes	87	Green	Yes
No	44	Blue	No
Yes	19	Red	No
No	32	Green	Yes
No	14	Blue	Yes



4 Yes, 2 No

Odds ratio (of yes): $4/2 = 2$

Odds ratio (of no) : $2/4 = 0.5$

$\text{Log}(2) = 0.7$

$\text{Log}(0.5) = -0.7$

1

First, make a “leaf (= guess for the y value)”

(initial leaf : log odds ratio of the highest value)

$\log(4/2) = 0.7$

(= first predicted y)



Put log odds ratio into
sigmoid function

$$\frac{1}{1 + e^{-\log(4/2)}} = 0.7$$

2. Gradient Boost Classification

Likes Popcorn	Age	Favorite Color	Loves Troll 2
Yes	12	Blue	Yes
Yes	87	Green	Yes
No	44	Blue	No
Yes	19	Red	No
No	32	Green	Yes
No	14	Blue	Yes



4 Yes, 2 No

Odds ratio (of yes): $4/2 = 2$

Odds ratio (of no) : $2/4 = 0.5$

$\text{Log}(2) = 0.7$

$\text{Log}(0.5) = -0.7$

Therefore, every sample will be classified 'Yes' for the initial guess!

1

First, make a "leaf (= guess for the y value)"

(initial leaf : log odds ratio of the highest value)

$\text{log}(4/2) = 0.7$ (= first predicted y)



Put log odds ratio into
sigmoid function

$$\frac{1}{1 + e^{-\text{log}(4/2)}} = 0.7$$



(threshold : 0.5)

Bigger than 0.5

-> classify into 'Yes'

Smaller than 0.5

-> classify into 'No'

2. Gradient Boost Classification

Likes Popcorn	Age	Favorite Color	Loves Troll 2	Residual
Yes	12	Blue	Yes	0.3
Yes	87	Green	Yes	0.3
No	44	Blue	No	-0.7
Yes	19	Red	No	-0.7
No	32	Green	Yes	0.3
No	14	Blue	Yes	0.3

(Yes = 1, No = 0)

(= 1(actual) – 0.7(predicted))

(= 0(actual) – 0.7(predicted))

2

Calculate Error based on the first prediction

2. Gradient Boost Classification

Likes Popcorn	Age	Favorite Color	Loves Troll 2	Residual
Yes	12	Blue	Yes	0.3
Yes	87	Green	Yes	0.3
No	44	Blue	No	-0.7
Yes	19	Red	No	-0.7
No	32	Green	Yes	0.3
No	14	Blue	Yes	0.3

(Yes = 1, No = 0)

(= 1(actual) – 0.7(predicted))

(= 0(actual) – 0.7(predicted))

Don't get confused!
0.7 is not "log odds ratio"!
It is "**probability**"

2

Calculate Error based on the first prediction

2. Gradient Boost Classification


Likes Popcorn	Age	Favorite Color	Loves Troll 2	Residual
Yes	12	Blue	Yes	0.3
Yes	87	Green	Yes	0.3
No	44	Blue	No	-0.7
Yes	19	Red	No	-0.7
No	32	Green	Yes	0.3
No	14	Blue	Yes	0.3

(Yes = 1, No = 0)

(= 1(actual) – 0.7(predicted))

(= 0(actual) – 0.7(predicted))

Don't get confused!
0.7 is not "log odds ratio"!
It is "**probability**"



2

Calculate **Error** based on the first prediction

3

Make a tree, **predicting the "residuals"**

2. Gradient Boost Classification

(in Regression)



A bit different in Classification



2. Gradient Boost Classification

(in Regression)

Average Weight

71.2

+

$lr \times$

Genders=F

Height<1.6

-14.7

4.8

Color not Blue

3.8

16.8

Height (m)	Favorite Color	Gender	Weight (kg)
1.6	Blue	Male	72.9

(If learning rate = 0.1)

$71.2 + (0.1 \times 16.8) = 72.9$




A bit different in Classification



(in Classification)

Average Weight

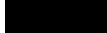


+

$lr \times$

Genders=F

Height<1.6



4.8

Color not Blue

3.8

16.8

Height (m)	Favorite Color	Gender	Weight (kg)
1.6	Blue	Male	72.9

(If learning rate = 0.1)

$71.2 + (0.1 \times 16.8) = 72.9$


Log odds ratio

Probability

MISMATCH! Can't just add them together!
How to make 'probability' compatible with 'log odds ratio'?

$$\frac{\sum \text{Residual}_i}{\sum [\text{Previous Probability}_i \times (1 - \text{Previous Probability}_i)]}$$

2. Gradient Boost Classification




With this formula, we can change the value of terminal node (which is **probability**) to be expressed in **log odds ratio**!

$$\frac{\sum \text{Residual}_i}{\sum [\text{Previous Probability}_i \times (1 - \text{Previous Probability}_i)]}$$

$$\frac{-0.7}{0.7 \times (1 - 0.7)} \quad (= -3.3)$$

$$\frac{0.3 + -0.7}{(0.7 \times (1 - 0.7)) + (0.7 \times (1 - 0.7))} \quad (= -1)$$

2. Gradient Boost Classification



Replace the value!
(to be compatible with 'log odds ratio')


$$\frac{\sum \text{Residual}_i}{\sum [\text{Previous Probability}_i \times (1 - \text{Previous Probability}_i)]}$$

$$\frac{-0.7}{0.7 \times (1 - 0.7)} \quad (= -3.3)$$

$$\frac{0.3 + -0.7}{(0.7 \times (1 - 0.7)) + (0.7 \times (1 - 0.7))} \quad (= -1)$$

2. Gradient Boost Classification

Likes Popcorn	Age	Favorite Color	Loves Troll 2
Yes	12	Blue	Yes
Yes	87	Green	Yes
No	44	Blue	No
Yes	19	Red	No



Prediction (before) : **0.7**

Prediction (after) : **$0.7 + (0.8 \times 1.4) = 1.82$**

convert **1.82** (log odds ratio) into **0.9** (probability)

Make a new prediction! (considering the tree newly made)



2. Gradient Boost Classification

Likes Popcorn	Age	Favorite Color	Loves Troll 2	Predicted Prob.	Residual
Yes	12	Blue	Yes	0.9	0.1
Yes	87	Green	Yes	0.5	0.5
No	44	Blue	No	0.5	-0.5
Yes	19	Red	No	0.1	-0.1
No	32	Green	Yes	0.9	0.1
No	14	Blue	Yes	0.9	0.1

(= 1(actual) – 0.9(predicted))

(= 1(actual) – 0.5(predicted))

4

Calculate error, based on the second prediction!

2-b. Gradient Boost

2. Gradient Boost Classification

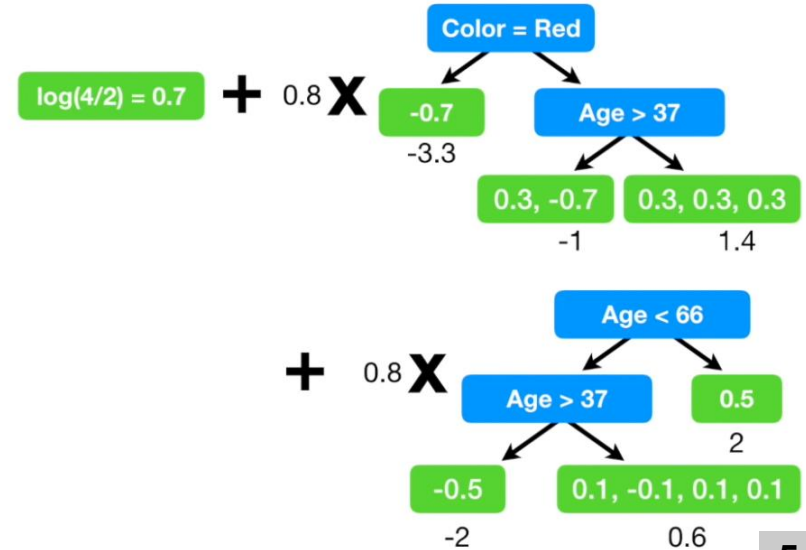
Likes Popcorn	Age	Favorite Color	Loves Troll 2	Predicted Prob.	Residual
Yes	12	Blue	Yes	0.9	0.1
Yes	87	Green	Yes	0.5	0.5
No	44	Blue	No	0.5	-0.5
Yes	19	Red	No	0.1	-0.1
No	32	Green	Yes	0.9	0.1
No	14	Blue	Yes	0.9	0.1

(= 1(actual) – 0.9(predicted))

(= 1(actual) – 0.5(predicted))

4

Calculate error, based on the second prediction!



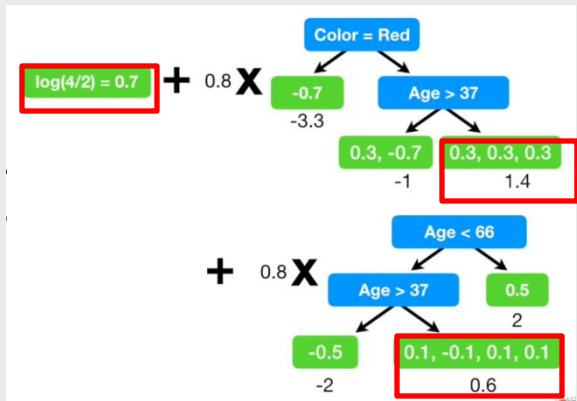
5

Keep making **sequential trees**!

2. Gradient Boost Classification

Ex) How would the person like below be classified?

Likes Popcorn	Age	Favorite Color	Loves Troll 2
Yes	25	Green	???



$$0.7 + (0.8 \times 1.4) + (0.8 \times 0.6) = 2.3$$

(into Logistic function)

$$f(2.3) = \exp(2.3) / (1 + \exp(2.3)) = 0.9$$

Therefore, it will be classified as 'Yes'

2. Gradient Boost Classification

Mathematical Expression

$$\text{Loss Function : } \text{LogLoss} = -\frac{1}{n} \sum_{i=0}^n [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

(also called **binary-cross entropy**)



2. Gradient Boost Classification


Mathematical Expression

$$\text{Loss Function : } \text{LogLoss} = -\frac{1}{n} \sum_{i=0}^n [y_i \log(\hat{y}_i) + (1 - y_i) \log(1 - \hat{y}_i)]$$

(also called **binary-cross entropy**)

$$- \left[\text{Observed} \times \log(p) + (1 - \text{Observed}) \times \log(1 - p) \right]$$

Case 1) Actual(=Observed) = 1, Predicted = 1



Log loss : **0**

Case 2) Actual(=Observed) = 1, Predicted = 0.5



Log loss : **0.69**

Case 3) Actual(=Observed) = 1, Predicted = 0



Log loss : **infinity**

2. Gradient Boost Classification

Gradient Boosted Regression Tree

1. **Input:** $D = \{(x_1, y_1), \dots, (x_N, y_N)\}, \theta, \gamma$
 2. **Output:** $F(x) = \sum_{i=0}^M F_i(x)$
 3. Initialize $F_0(x) = \arg \min_{\beta} \sum_{i=0}^N L(y_i, \beta)$
 4. **While** ($m < M$)
 5. $d_i = -[\partial L(y_i, F(x_i)) / \partial F(x_i)]_{F(x_i)=F_{m-1}(x_i)}$
 6. $\vartheta = \{(x_i, d_i)\}, i = 1, N$
 7. $g(x) = \text{FITREGRETREE}(\vartheta, \theta)$
 8. $\rho_m = \arg \min_{\rho} \sum_{i=1}^N L(y_i, F_{m-1}(x) + \rho g(x))$
 9. $F_m(x) = F_{m-1}(x) + \gamma \rho_m g(x)$
 13. **End while**
-

2. Gradient Boost Classification

Gradient Boosted	Classification	Tree
1.	Input: $D = \{(x_1, y_1), \dots, (x_N, y_N)\}, \theta, \gamma$	
2.	Output: $F(x) = \sum_{i=0}^M F_i(x)$	
3.	Initialize $F_0(x) = \arg \min_{\beta} \sum_{i=0}^N L(y_i, \beta)$	
4.	While ($m < M$)	
5.	$d_i = -[\partial L(y_i, F(x_i)) / \partial F(x_i)]_{F(x_i) = F_{m-1}(x_i)}$	
6.	$\vartheta = \{(x_i, d_i)\}, i = 1, N$	
7.	$g(x) = \text{FITREGRETREE}(\vartheta, \theta)$	
8.	$\rho_m = \arg \min_{\rho} \sum_{i=1}^N L(y_i, F_{m-1}(x) + \rho g(x))$	
9.	$F_m(x) = F_{m-1}(x) + \gamma \rho_m g(x)$	
13.	End while	

Almost the same as the Regression Model!

2-b. Gradient Boost

2. Gradient Boost Classification

Gradient Boosted Regression Tree

1. **Input:** $D = \{(x_1, y_1), \dots, (x_N, y_N)\}, \theta, \gamma$
2. **Output:** $F(x) = \sum_{i=0}^M F_i(x)$
3. Initialize $F_0(x) = \arg \min_{\beta} \sum_{i=0}^N L(y_i, \beta)$
4. **While** ($m < M$)
 - Loss Function : **MSE**
 - $d_i = -[\partial L(y_i, F(x_i)) / \partial F(x_i)]_{F(x_i)=F_{m-1}(x_i)}$
 - $\vartheta = \{(x_i, d_i)\}, i = 1, N$
 - $g(x) = \text{FITREGTREETREE}(\vartheta, \theta)$
 - $\rho_m = \arg \min_{\rho} \sum_{i=1}^N L(y_i, F_{m-1}(x) + \rho g(x))$
 - $F_m(x) = F_{m-1}(x) + \gamma \rho_m g(x)$
13. **End while**

$$y = f(x) + e1$$
$$e1 = g(x) + e2$$
$$e2 = h(x) + e3$$

$$Y = f(x) + g(x) + h(x) + \dots$$

First prediction = “**average**” of all y values

Height (m)	Color	Gender	Weight (kg)	Residual
1.6	Blue	Male	80	16.8
1.6	Green	Female	76	4.8
1.5	Blue	Female	66	-15.2
1.8	Red	Male	73	1.8
1.5	Green	Male	77	5.8
1.4	Blue	Female	57	-14.2



$$Y = f(x) + g(x) + e2$$

2. Gradient Boost Classification

Gradient Boosted Classification Tree

1. **Input:** $D = \{(x_1, y_1), \dots, (x_N, y_N)\}, \theta, \gamma$
2. **Output:** $F(x) = \sum_{i=0}^M F_i(x)$
3. Initialize $F_0(x) = \arg \min_{\beta} \sum_{i=0}^N L(y_i, \beta)$
4. **While** ($m < M$)

Loss Function : **Log Loss**
5. $d_i = -[\partial L(y_i, F(x_i)) / \partial F(x_i)]_{F(x_i)=F_{m-1}(x_i)}$
6. $\vartheta = \{(x_i, d_i)\}, i = 1, N$
7. $g(x) = \text{FITREGTREETREE}(\vartheta, \theta)$
8. $\rho_m = \arg \min_{\rho} \sum_{i=1}^N L(y_i, F_{m-1}(x) + \rho g(x))$
9. $F_m(x) = F_{m-1}(x) + \gamma \rho_m g(x)$
13. **End while**

$$y = f(x) + e1$$

$$e1 = g(x) + e2$$

$$e2 = h(x) + e3$$

$$Y = f(x) + g(x) + h(x) + \dots$$

First prediction = “log odds ratio” of the mode

Height (m)	Color	Gender	Weight (kg)	Residual
1.6	Blue	Male	88	16.8
1.6	Green	Female	76	4.8
1.5	Blue	Female	96	-15.2
1.8	Red	Male	73	1.8
1.5	Green	Male	77	5.8
1.4	Blue	Female	57	-14.2



$$Y = f(x) + g(x) + e2$$



2-c. XGBoost



1. Characteristics of XGBoost

XGBoost (eXtreme Gradient Boost)

Consists of lots of parts!

- 1) Gradient Boost
- 2) Regularization
- 3) Unique Regression Tree
- 4) Approximate Greedy Algorithm
- 5) Weighted Quantile Sketch
- 6) Sparsity-Aware Split Finding
- 7) Parallel Learning
- 8) Cache-Aware Access
- 9) Blocks for Out-of-Core Computation

1. XGBoost Regression

2. XGBoost Classification



2. XGBoost Regression





2. XGBoost Regression

Unique Regression Tree!

Example with simple data)

X : drug dosage


Y : drug Effectiveness



1

Initial Prediction : **0.5 (default)**
(both in Regression & Classification)

2. XGBoost Regression




(same as non-extreme GB)

2

Then, predict the “residual”

(build the next model (tree) predicting the residuals)


2. XGBoost Regression



(same as non-extreme GB) **2**

Then, predict the “**residual**”

(build the next model (tree) predicting the residuals)



HOW?

Step 1.

Put all the residuals to the same leaf

-10.5, 6.5, 7.5, -7.5

And calculate a “**similarity score**”

$$\text{Similarity Score} = \frac{\text{Sum of Residuals, Squared}}{\text{Number of Residuals} + \lambda}$$

(lambda : regularization term)


$$\frac{(-10.5 + 6.5 + 7.5 + -7.5)^2}{4 + \mathbf{0}} = 4$$

2. XGBoost Regression

HOW?

Step 2.

Split the leaf, which gives the most “Gain”



1 Drug Dosage < 15

2 Drug Dosage < 22.5

3 Drug Dosage < 15


Which split criterion
should we choose?

2. XGBoost Regression

HOW?


Step 2.

Split the leaf, which gives the most **“Gain”**



- 1 Drug Dosage < 15
- 2 Drug Dosage < 22.5
- 3 Drug Dosage < 15

Which split criterion
should we choose?



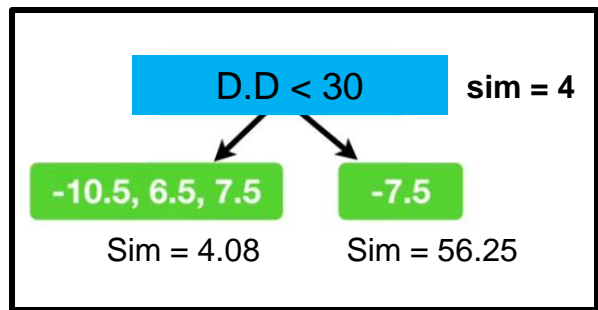


2. XGBoost Regression

- 1 Drug Dosage < 15
- 2 Drug Dosage < 22.5
- 3 Drug Dosage < 30

Which split criterion
should we choose?

Example)

$$14.08 = \frac{(6.5 + 7.5 + -7.5)^2}{3 + 0}$$




2. XGBoost Regression

- 1 Drug Dosage < 15
- 2 Drug Dosage < 22.5
- 3 Drug Dosage < 30

Which split criterion
should we choose?


Example)

$$14.08 = \frac{(6.5 + 7.5 + -7.5)^2}{3 + 0}$$




Gain

$$(110.25 + 14.08) - 4 = 120.33$$



Gain

$$(8 + 0) - 4 = 4$$



Gain

$$(4.08 + 56.25) - 4 = 56.33$$


2. XGBoost Regression

- 1 Drug Dosage < 15
- 2 Drug Dosage < 22.5
- 3 Drug Dosage < 30

Which split criterion
should we choose?


Example)

$$14.08 = \frac{(6.5 + 7.5 + -7.5)^2}{3 + 0}$$



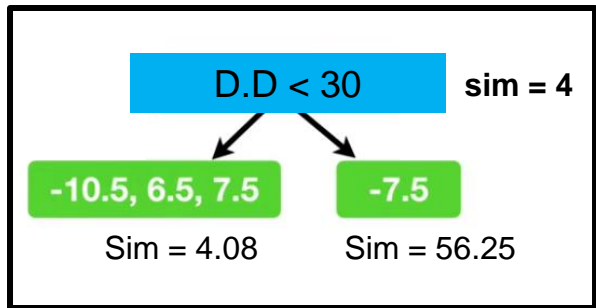
Gain

$$(110.25 + 14.08) - 4 = 120.33$$



Gain

$$(8 + 0) - 4 = 4$$

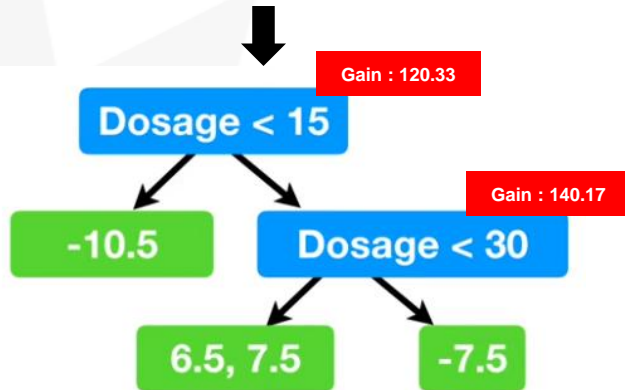



Gain

$$(4.08 + 56.25) - 4 = 56.33$$

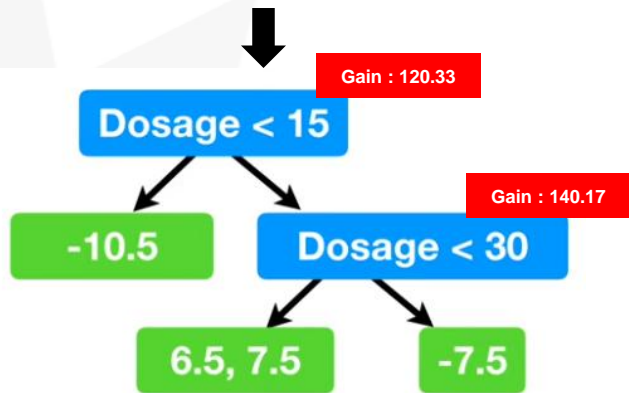
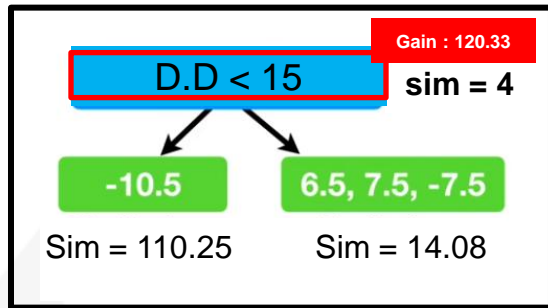
2. XGBoost Regression

choose “D.D<15” as the criterion



2. XGBoost Regression

choose "D.D < 15" as the criterion




HOW?

Step 3.

Prune the tree (if necessary)

To Prevent
Overfitting!



New parameter : '**gamma**'

NO GAIN,
NO PRUNE

If $\gamma = 130$:

($140.17 > 130$) -> Prune X

If $\gamma = 150$:

($140.17 < 150$) -> Prune O -> ($120.33 < 150$) -> Prune O

2. XGBoost Regression

Another method for preventing overfitting (regularization)

$$\text{Similarity Score} = \frac{\text{Sum of Residuals, Squared}}{\text{Number of Residuals} + \lambda}$$

2. XGBoost Regression

Another method for preventing overfitting (regularization)

$$\text{Similarity Score} = \frac{\text{Sum of Residuals, Squared}}{\text{Number of Residuals} + \lambda}$$

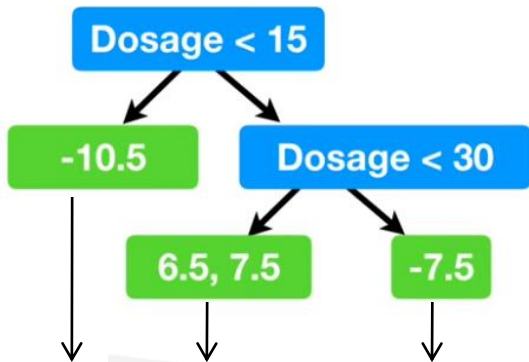
What happens if **lambda** gets larger?

lambda \uparrow -> similarity score \downarrow -> **More chance of getting pruned**

2. XGBoost Regression

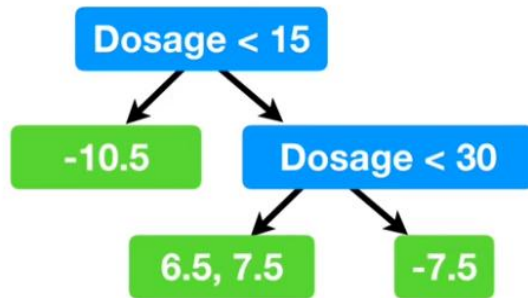
Step 4.

Predict output values of each leaf !



these are all RESIDUALS, not output values!

$$\text{Output Value} = \frac{\text{Sum of Residuals}}{\text{Number of Residuals} + \lambda}$$




How will these values (of leaves) change?

2. XGBoost Regression


Step 4.

Predict output values of each leaf !



these are all RESIDUALS, not output values!

$$\text{Output Value} = \frac{\text{Sum of Residuals}}{\text{Number of Residuals} + \lambda}$$



2. XGBoost Regression


Step 5.

Predict the final outcome!

Predicted Drug
Effectiveness
0.5

+

Learning Rate **X**



...

...

...




2. XGBoost Regression

Step 5.

Predict the final outcome!

Getting a more precise prediction as the model is made!

(But **be careful of overfitting!**)





3. XGBoost Classification






3. XGBoost Classification

Example with simple data)

X : drug dosage


Y : Effective / Not Effective



1

Initial Prediction : 0.5 (default)
(both in Regression & Classification)

3. XGBoost Classification




(same as non-extreme GB)

2

Then, predict the “residual”

(build the next model (tree) predicting the residuals)

3. XGBoost Classification



(same as non-extreme GB) **2**

Then, predict the “**residual**”

(build the next model (tree) predicting the residuals)

HOW?

Step 1.

Put all the residuals to the same leaf

-0.5, 0.5, 0.5, -0.5

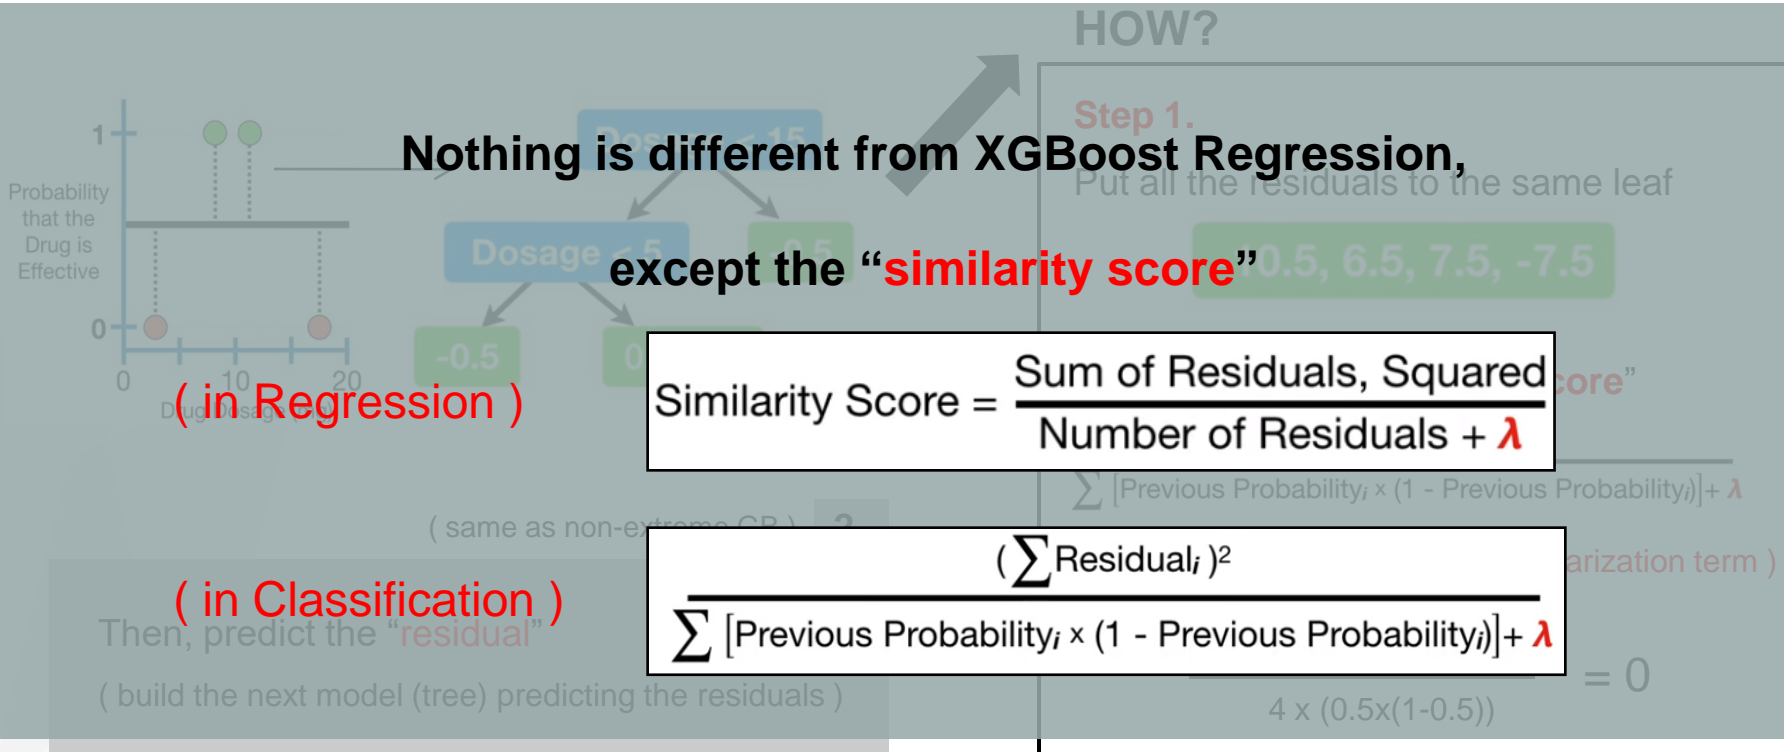
And calculate a “**similarity score**”

$$\frac{(\sum \text{Residual}_i)^2}{\sum [\text{Previous Probability}_i \times (1 - \text{Previous Probability}_i)] + \lambda}$$

(lambda : regularization term)

$$\frac{(-0.5 + 0.5 + 0.5 + -0.5)^2}{4 \times (0.5 \times (1 - 0.5))} = 0$$

3. XGBoost Classification



2. XGBoost Regression

HOW?

Step 2.

Split the leaf, which gives the most “Gain”



1 Drug Dosage < 15

2 Drug Dosage < 22.5

3 Drug Dosage < 15

Which split criterion
should we choose?


2-c. XGBoost

2. XGBoost Regression

HOW?

Step 2.

Split the leaf, which gives the most **“Gain”**




1 Drug Dosage < 15

2 Drug Dosage < 22.5

3 Drug Dosage < 15

Which split criterion
should we choose?

Calculate the
similarity score!



Repeat this
for the next leaf



3. XGBoost Classification

Step 3.

Pruning the tree

- parameter '**gamma**'
(minimum gain to make a split!)
- parameter '**lambda**'
(regularization term)



3. XGBoost Classification

Step 3.

Pruning the tree

- parameter '**gamma**'
(minimum gain to make a split!)
- parameter '**lambda**'
(regularization term)
- **minimum number of residuals in each leaf**
(new term, '**cover**')

3. XGBoost Classification

Step 3.

Pruning the tree

- parameter '**gamma**'
(minimum gain to make a split!)
- parameter '**lambda**'
(regularization term)
- **minimum number of residuals in each leaf**
(new term, '**cover**')

$$(\sum \text{Residual}_i)^2$$

$$\sum [\text{Previous Probability}_i \times (1 - \text{Previous Probability}_i)] + \lambda$$

(in regression)

$$\text{Similarity Score} = \frac{\text{Sum of Residuals, Squared}}{\text{Number of Residuals} + \lambda}$$

- cover can be defined as 'number of residuals'
- default value : 1 -> means that it has no affect on the growth of tree

3. XGBoost Classification

Step 3.

Pruning the tree

- parameter '**gamma**'
(minimum gain to make a split!)
- parameter '**lambda**'
(regularization term)
- minimum number of residuals in each leaf
(new term, '**cover**')

$$(\sum \text{Residual}_i)^2$$

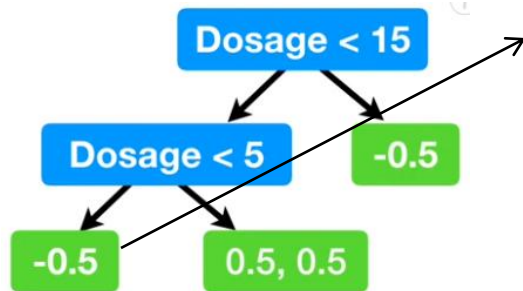
$$\sum [\text{Previous Probability}_i \times (1 - \text{Previous Probability}_i)] + \lambda$$

(in regression)

$$\text{Similarity Score} = \frac{\text{Sum of Residuals, Squared}}{\text{Number of Residuals} + \lambda}$$

- cover can be defined as 'number of residuals'
- default value : 1 -> means that it has no affect on the growth of tree

Example)



'cover' for this leaf is
 $1 \times (0.5 \times (1-0.5)) = 0.25$

3. XGBoost Classification

Step 3.

Pruning the tree

- parameter '**gamma**'
(minimum gain to make a split!)
- parameter '**lambda**'
(regularization term)
- minimum number of residuals in each leaf
(new term, '**cover**')

$$(\sum \text{Residual}_i)^2$$


$$\sum [\text{Previous Probability}_i \times (1 - \text{Previous Probability}_i)] + \lambda$$

(in regression)

$$\text{Similarity Score} = \frac{\text{Sum of Residuals, Squared}}{\text{Number of Residuals} + \lambda}$$

- cover can be defined as 'number of residuals'
- default value : 1 -> means that it has no affect on the growth of tree

Example)




we call this parameter 'cover' as "**min_child_weight**"

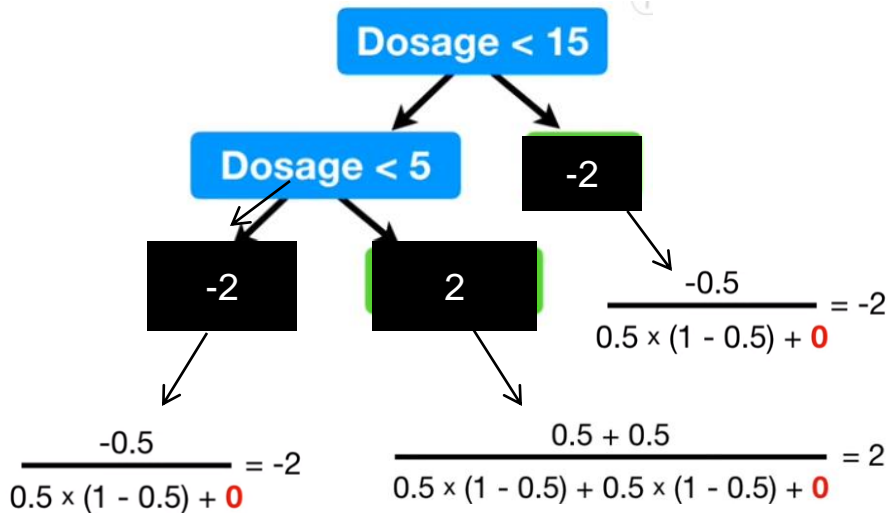
3. XGBoost Classification

Step 4.

Predict the output values of each leaf



$$\text{Output value} = \frac{(\sum \text{Residual}_i)}{\sum [\text{Previous Probability}_i \times (1 - \text{Previous Probability}_i)] + \lambda}$$



3. XGBoost Classification

Step 5.

Predict the final outcome

Predicted Drug
Effectiveness

0.5

+

learning rate

×

Dosage < 15

Dosage < 5

-0.5

Output
= -2

-0.5

Output
= -2

Output
= 2

0.5, 0.5

Probability

Log Odds Ratio

...

...

...

3. XGBoost Classification

Step 5.

Predict the final outcome

Predicted Drug
Effectiveness

0.5



learning rate



Dosage < 15

Dosage < 5

-0.5

Output
= -2

-0.5

Output
= -2

0.5, 0.5

Output
= 2

...

...

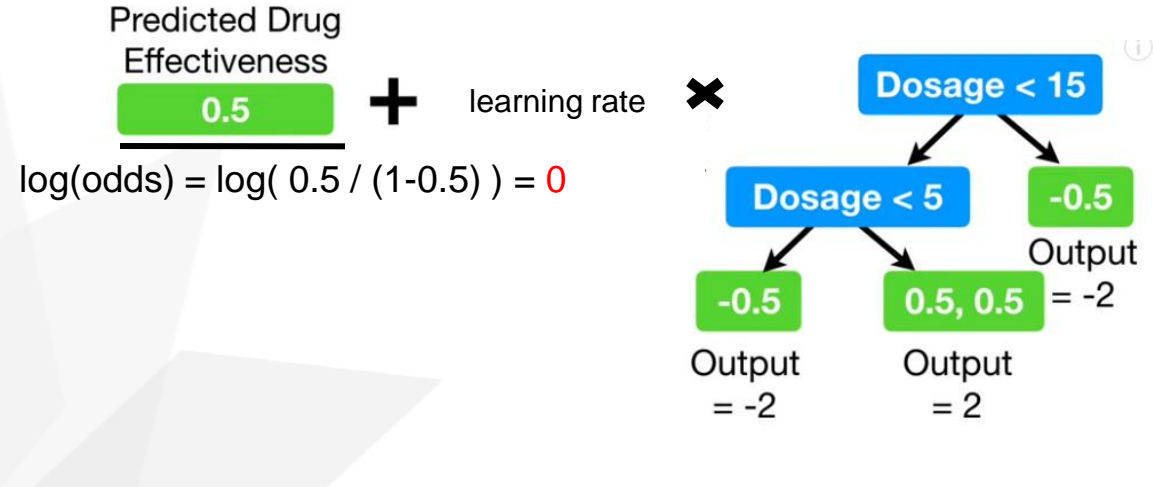
...

0.5 is “probability”, so we have to
convert this into “log odds ratio”

3. XGBoost Classification


Step 5.

Predict the final outcome



3. XGBoost Classification

Ex) What is the predicted class for the sample below, with the red color? (let learning rate = 0.3)



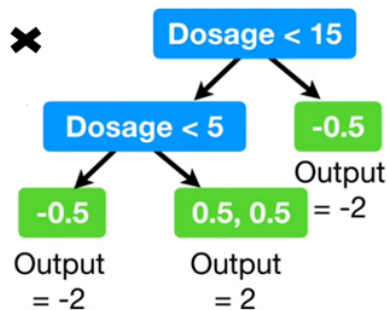
Predicted Drug Effectiveness

$$\frac{0.5}{\log(\text{odds}) = \log(0.5 / (1-0.5)) = 0}$$

+

learning rate

×



Answer :

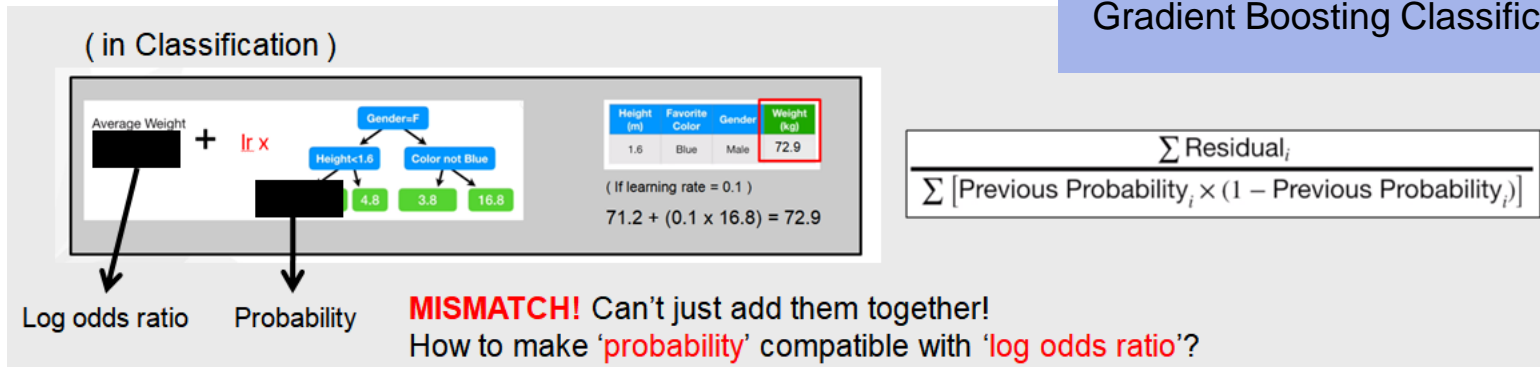
$$1) 0 + (0.3 \times (-2)) = -0.6$$

(this is in terms of 'log(odds)' ! Have to convert into probability)

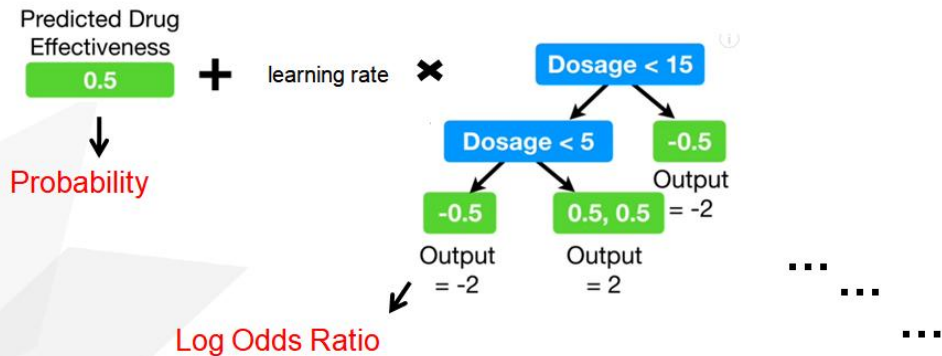
$$2) \text{Probability} = \frac{e^{-0.6}}{1 + e^{-0.6}} = 0.35 \rightarrow \text{Predicted as 'NO EFFECT'}$$

3. XGBoost Classification

Gradient Boosting Classification



XGBoost Classification



3. XGBoost Classification

Parameters of XGBoost

Parameter	Description	Default Value
max_depth	controls the maximum depth of each tree (used to control over-fitting)	6
subsample	specifies the fraction of observations to be randomly sampled at each tree (adds randomness)	1
eta	the learning rate	0.3
nrounds	the number of trees to be produced (similar to ntree)	100–1000
gamma	controls the minimum loss reduction required to make a node split (used to control over-fitting)	0
min_child_weight	Specifies the minimum sum of instance weight of all the observations required in a child (used to control over-fitting)	1
colsample_bytree	Specifies the number of features to consider when searching for the best node split (adds randomness)	1

3. XGBoost Classification

Mathematical Expressions

<https://brunch.co.kr/@snobberys/137>

- Model: assuming we have K trees

$$\hat{y}_i = \sum_{k=1}^K f_k(x_i), \quad f_k \in \mathcal{F}$$

- Objective

$$Obj = \sum_{i=1}^n l(y_i, \hat{y}_i) + \sum_{k=1}^K \Omega(f_k)$$

Training loss

Complexity of the Trees

3. XGBoost Classification

Mathematical Expressions

- **Goal** $Obj^{(t)} = \sum_{i=1}^n l(y_i, \hat{y}_i^{(t-1)} + f_t(x_i)) + \Omega(f_t) + constant$
 - Seems still complicated except for the case of square loss
- Take Taylor expansion of the objective
 - Recall $f(x + \Delta x) \simeq f(x) + f'(x)\Delta x + \frac{1}{2}f''(x)\Delta x^2$
 - Define $g_i = \partial_{\hat{y}^{(t-1)}} l(y_i, \hat{y}^{(t-1)})$, $h_i = \partial_{\hat{y}^{(t-1)}}^2 l(y_i, \hat{y}^{(t-1)})$

$$Obj^{(t)} \simeq \sum_{i=1}^n \left[l(y_i, \hat{y}_i^{(t-1)}) + g_i f_t(x_i) + \frac{1}{2} h_i f_t^2(x_i) \right] + \Omega(f_t) + constant$$



3. XGBoost Classification

Mathematical Expressions

$$\Omega(f_t) = \underbrace{\gamma T}_{\text{Number of leaves}} + \underbrace{\frac{1}{2}\lambda \sum_{j=1}^T w_j^2}_{\text{L2 norm of leaf scores}}$$

$$Obj^{(t)} \simeq \sum_{i=1}^n \left[l(y_i, \hat{y}_i^{(t-1)}) + g_i f_t(x_i) + \frac{1}{2} h_i f_t^2(x_i) \right] + \underbrace{\Omega(f_t)}_{\uparrow} + constant$$



At time 't'

$$\sum_{i=1}^n \left[g_i f_t(x_i) + \frac{1}{2} h_i f_t^2(x_i) \right] + \Omega(f_t)$$

- where $g_i = \partial_{\hat{y}^{(t-1)}} l(y_i, \hat{y}^{(t-1)})$, $h_i = \partial_{\hat{y}^{(t-1)}}^2 l(y_i, \hat{y}^{(t-1)})$

3. XGBoost Classification

Mathematical Expressions

- Define the instance set in leaf j as $I_j = \{i | q(x_i) = j\}$
- Regroup the objective by each leaf

$$\begin{aligned} Obj^{(t)} &\simeq \sum_{i=1}^n \left[g_i f_t(x_i) + \frac{1}{2} h_i f_t^2(x_i) \right] + \Omega(f_t) \\ &= \sum_{i=1}^n \left[g_i w_{q(x_i)} + \frac{1}{2} h_i w_{q(x_i)}^2 \right] + \gamma T + \lambda \frac{1}{2} \sum_{j=1}^T w_j^2 \\ &= \sum_{j=1}^T \left[\left(\sum_{i \in I_j} g_i \right) w_j + \frac{1}{2} \left(\sum_{i \in I_j} h_i + \lambda \right) w_j^2 \right] + \gamma T \end{aligned}$$

- This is sum of T independent quadratic functions

3. XGBoost Classification

Mathematical Expressions

G : sum of g

H : sum of h

$$\operatorname{argmin}_x Gx + \frac{1}{2}Hx^2 = -\frac{G}{H}, \quad H > 0 \quad \min_x Gx + \frac{1}{2}Hx^2 = -\frac{1}{2}\frac{G^2}{H}$$

$$w_j^* = -\frac{G_j}{H_j + \lambda} \quad Obj = -\frac{1}{2} \sum_{j=1}^T \frac{G_j^2}{H_j + \lambda} + \gamma T$$

↑
This measures how good a tree structure is!

Make a tree that makes the best split according to the following gain!

$$Gain = \frac{G_L^2}{H_L + \lambda} + \frac{G_R^2}{H_R + \lambda} - \frac{(G_L + G_R)^2}{H_L + H_R + \lambda} - \gamma$$



3. Boosting with Python

- 1) Adaboost
- 2) Gradient Boosting
- 3) XGBoost



4. Homework

Homework



HW1) **Classification**


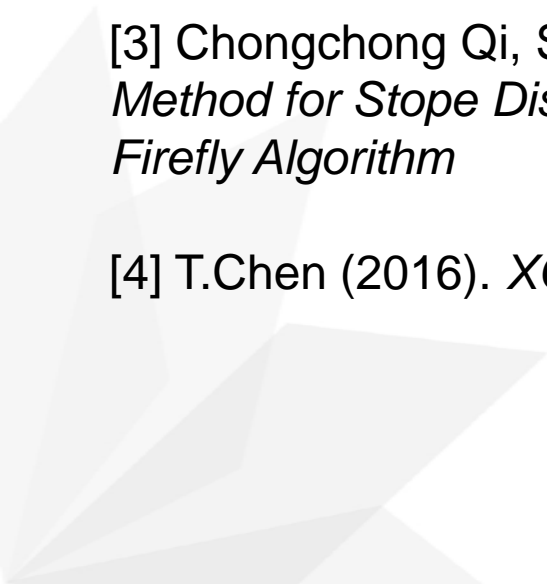
HW2) **Regression**

- Choose any two of the ten datasets below

<https://machinelearningmastery.com/standard-machine-learning-datasets/>

Make a prediction using the models below (5-fold Cross Validation)
(Gradient Boosting, XGBoost, LightGBM + Stacking (next week))

Reference

- 
- [1] Josh Starmer(2019). *StatQuest with Josh Starmer*
 - [2] Seoul National University, Hyeong In Choi (2017). *Lecture 11: Boosting (I) AdaBoost (Draft: version 0.9.2)*
 - [3] Chongchong Qi, S.M.ASCE, Andy Fourie, Xu Zhao (2018). *Back-Analysis Method for Slope Displacements Using Gradient-Boosted Regression Tree and Firefly Algorithm*
 - [4] T.Chen (2016). *XGBoost : A Scalable Tree Boosting System*
- 



Thank you